

NATURAL LANGUAGE PROCESSING (PRACTICE)

NLP 251 - Lab 3: Language Models



Department of Computer Science and Engineering
Ho Chi Minh University of Technology, VNU-HCM

Language Model

Language Model

Language Model: A language model is a model that predicts the probability of a sequence of words.

$$P(W) = P(w_1, w_2, \dots, w_n)$$

Example

- $S_1 = \text{"The cat jumped over the dog."} \Rightarrow P(S_1) \approx 1$
- $S_2 = \text{"The jumped cat the over dog."} \Rightarrow P(S_2) \approx 0$

Application

- Machine Translation
 - $P(\text{high winds tonight}) > P(\text{large winds tonight})$
- Text Correction
 - The office is about fifteen *minutes* from my house
 - $P(\text{"about fifteen minutes from"}) > P(\text{about fifteen minuets from})$
- Speech Recognition
 - $P(\text{I saw a van}) > P(\text{eyes awe of an})$
- Handwriting Recognition
 - $P(\text{Act naturally}) > P(\text{Abt naturally})$
- Summarization, Q&A, etc.

Conditional probability

Formula

$$P(A|B) = \frac{P(A \cap B)}{P(B)} \text{ or } P(A, B) = P(A) \cdot P(B|A)$$

$$P(A, B, C, D) = P(A) \cdot P(B|A) \cdot P(C|A, B) \cdot P(D|A, B, C)$$

Therefore

$$P(S) = P(w_1) \cdot P(w_2|w_1) \cdot P(w_3|w_1, w_2) \cdots P(w_n|w_1, w_2, w_3, \dots, w_n)$$

Example

$$\begin{aligned} P(\text{Computer, can, recognize, speech}) &= P(\text{Computer}) \cdot \\ &P(\text{can}|\text{Computer}) \cdot \\ &P(\text{recognize}|\text{Computer can}) \cdot \\ &P(\text{speech}|\text{Computer can recognize}) \end{aligned}$$

N-gram Model

Markov hypothesis

Markov property

giả định rằng xác suất của từ hiện tại chỉ phụ thuộc vào từ trước đó

$$P(S) = \prod_{i=1}^n P(w_i | w_1, w_2, \dots, w_{i-1}) \Rightarrow P(S) = \prod_{i=1}^n P(w_i | w_{i-1})$$

Example:

$$\begin{aligned} P(\text{Computer, can, recognize, speech}) &= P(\text{Computer}) \cdot \\ &\quad P(\text{can} | \text{Computer}) \cdot \\ &\quad P(\text{recognize} | \text{Computer can}) \cdot \\ &\quad P(\text{speech} | \text{Computer can recognize}) \\ P(\text{Computer, can, recognize, speech}) &= P(\text{Computer}) \cdot P(\text{can} | \text{Computer}) \cdot \\ &\quad P(\text{recognize} | \text{can}) \cdot P(\text{speech} | \text{recognize}) \end{aligned}$$

N-GRAM model

không phụ thuộc vào từ trước đó

- **Unigram (1-gram):** probability of a word independent of previous words

$$P(w_1 w_2 \dots w_n) \approx \prod_i P(w_i)$$

chỉ phụ thuộc vào từ trước đó

- **Bigram (2-gram):** probability of a word given the previous word

$$P(w_1 w_2 \dots w_n) \approx \prod_i P(w_i | w_{i-1})$$

phụ thuộc vào 2 từ trước đó

- **Trigram (3-gram):** probability of a word dependent on two previous words

$$P(w_1 w_2 \dots w_n) \approx \prod_i P(w_i | w_{i-1}, w_{i-2})$$

phụ thuộc vào N từ trước đó

- **N-gram:** probability of a word dependent on N previous words

$$P(w_1 w_2 \dots w_n) \approx \prod_i P(w_i | w_{i-1}, w_{i-2}, w_{i-N})$$

Likelihood Estimate

Likelihood Estimate of Bigram (2-gram):

$$P(w_i | w_{i-1}) = \frac{\text{count}(w_{i-1}, w_i)}{\text{count}(w_{i-1})} \quad \begin{array}{l} \text{số lần cặp từ xuất hiện} \\ \text{số lần từ đó xuất hiện} \end{array}$$

Example:

- $\langle s \rangle$ I am Sam $\langle /s \rangle$
- $\langle s \rangle$ Sam I am $\langle /s \rangle$
- $\langle s \rangle$ I do not like green eggs and ham $\langle /s \rangle$

$$P(I|\langle s \rangle) = \frac{2}{3} = .67 \quad P(\text{Sam}|\langle s \rangle) = \frac{1}{3} = .33 \quad P(\text{am}|I) = \frac{2}{3} = .67$$

$$P(\langle /s \rangle|\text{Sam}) = \frac{1}{2} = 0.5 \quad P(\text{Sam}|\text{am}) = \frac{1}{2} = .5 \quad P(\text{do}|I) = \frac{1}{3} = .33$$

Smoothing and Zeros

The Zero-Probability Problem

Nếu mô hình chưa từng thấy 1 n-gram trong tập huấn luyện, nó sẽ gán luôn xác suất bằng 0 cho cả câu chứa n-gram đó

--> Không hợp lý

Training set:

- ... denied the allegations
- ... denied the reports
- ... denied the claims
- ... denied the request

Test set:

- ... denied the offer
- ... denied the loan
- ...
- ...

$P(\text{"offer"}|\text{denied the}) = 0$. (This means we will assign a probability of 0 to the above sentence)

Smoothing

- **Laplace smoothing:** +1 để tránh xác suất bằng 0

$$P_{\text{Laplace}}(w_i | w_{i-1}) = \frac{c(w_{i-1}, w_i) + 1}{c_{w_{i-1}} + V}$$

V: số unique words của tập từ vựng

- **Linear Interpolation:**

$$P(w_n | w_{n-1} w_{n-2}) = \lambda_1 P(w_n | w_{n-1} w_{n-2}) + \lambda_2 P(w_n | w_{n-1}) + \lambda_3 P(w_n)$$

cân bằng giữa độ chính xác (context dài) và độ phủ(context ngắn)

with $\lambda_1 + \lambda_2 + \lambda_3 = 1$. For interpolation with context-conditioned weights, where each lambda takes an argument that is the two prior word context:

$$P(w_n | w_{n-2} w_{n-1}) = \lambda_1(w_{n-2}, w_{n-1}) P(w_n) + \lambda_2(w_{n-2}, w_{n-1}) P(w_n | w_{n-1}) \\ + \lambda_3(w_{n-2}, w_{n-1}) P(w_n | w_{n-2} w_{n-1})$$

thêm điều kiện ==> mô hình thông minh hơn

Smoothing

Use valid set to choose parameter :

- Keep N-gram probabilities fixed (from training data)
- Then find λ that maximize probability on valid set:

$$\log P(w_1 \dots w_n | M(\lambda_1 \dots \lambda_k)) = \sum_i \log P_{M(\lambda_1 \dots \lambda_k)}(w_i | w_{i-1})$$

Chọn tập lambda này sao cho hàm log là lớn nhất

Besides the methods mentioned above, we also have methods such a

- **Good-Turing**
- **Witten-Bell**
- **Kneser-Ney**
- **Backoff**

Huge Language Models and Stupid Backoff

- Solving large-value problems, e.g., Google N-gram corpus
- Pruning
 - Only store N-grams with frequency $>$ threshold
 - Remove higher-order n-gram entries
- Efficiency
 - Use efficient data structures like tries
 - Bloom filters: approximate language model matching
 - Store words as indices, not strings
 - Use Huffman coding to convert large words into 2 bytes
 - Probability quantization (4-8 bits instead of 8-byte float)

Stupid backoff algorithm helps the model maintain a manageable size while still achieving reasonable effectiveness.

Model Evaluation

Perplexity Metric

hàm ngược của hàm xác suất

Perplexity is the inverse probability of the test set, normalized by the number of words
 N là số token trong câu đang xét

$$PP(W) = P(w_1 w_2 \dots w_N)^{\frac{1}{N}} = \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}} = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_1 \dots w_{i-1})}}$$

- Minimizing perplexity is equivalent to maximizing probability
- Low perplexity = good model

With **WSJ dataset**: training set of 38 million words, test set of 1.5 million words

N-gram Order	Unigram	Bigram	Trigram
Perplexity	962	170	109

Shannon Visualization Method

Algorithm

- Choose a random bigram ($\langle s \rangle$, w) according to its probability
- Now choose a random bigram (w , x) according to its probability
- And so on until we choose $\langle /s \rangle$
- Then string the words together

ký hiệu bắt đầu
 $\langle s \rangle$ I
 I want
 want to
 to eat
 eat Chinese
 Chinese food ký hiệu kết thúc
 food $\langle /s \rangle$
 I want to eat Chinese food

N-Value

số lượng n-gram = V^n
 - V là số unique words trong vocab
 - n là số vocab xét trong 1 gram

What is the appropriate value of n ?

- Theoretically, very difficult to determine
- However: as large as possible (\rightarrow approaches the "perfect" model)
- Empirically, $n = 3$ is common
 - Parameter estimation? (confidence, data, storage, space, ...)
 - 4 is too large: $|V| = 60k \rightarrow 1.296 \times 10^{19}$ parameters
 - However: 6 – 7 possible with sufficient data: in practice, we can recover from 7-grams!

THANKS FOR
LISTENING!